GlycoSMILES2BAP: An Automated Pipeline for Stereochemistry-Preserving Glycan Structure Prediction with AlphaFold 3

Qiang Xia
Zhejiang Xinghe Tea Technology Co., Ltd.
Hangzhou, Zhejiang, China
xiaqiang@xinghetea.com

Abstract

Motivation: AlphaFold 3 (AF3) has revolutionized protein structure prediction and expanded capabilities to include glycan ligands. However, recent work by Huang et al. (2025) identified a critical limitation: standard input formats (SMILES, userCCD) produce stereochemically incorrect glycan structures, while the stereochemistry-preserving CCD+bondedAtomPairs (BAP) format requires impractical manual atom-by-atom specification.
Results: We present GlycoSMILES2BAP, an automated pipeline that converts standard glycan notations (IUPAC-condensed, WURCS) to AF3-compatible CCD+BAP format. The pipeline integrates three core modules: (1) a CCD mapper supporting 28+ monosaccharide configurations with anomeric position tracking, (2) a topology parser extracting linkage information from branched structures, and (3) a BAP generator producing explicit atom-pair bond specifications. Validated on a curated benchmark of 50 diverse glycan structures, GlycoSMILES2BAP achieves 98.5% epimer accuracy (95% CI: [96.2%, 99.8%]), 98.2% anomeric accuracy (95% CI: [95.8%, 99.6%]), and 96.8% linkage accuracy (95% CI: [93.5%, 99.2%]) compared to ∼60% for SMILES-based approaches. Processing time is <1 second per structure versus 30–60 minutes for manual specification. Systematic ablation studies confirm that all three core modules contribute significantly to the high accuracy ($p < 0.001$). Furthermore, validation against 10 literature-reported structure errors demonstrated successful correction in all tested cases (n=10), and database-scale processing achieved 94% success rate on 100 GlyTouCan representative structures. Complexity analysis of the validation set revealed balanced coverage: 35% linear, 30% mono-branched, 25% bi-antennary, and 10% multi-antennary structures, with an average of 3.2 distinct monosaccharide types per structure.
Conclusions: GlycoSMILES2BAP eliminates the input format barrier for AF3 glycan modeling, enabling accurate, reproducible structure prediction for the glycobiology community. The tool successfully corrects common stereochemistry errors and scales to database-level processing, making it suitable for both individual structure validation and large-scale structural glycomics. The tool is open-source and readily integrable with existing glycan databases (GlyTouCan, GlyGen).
Availability: Open-source implementation available at `https://github.com/xiaqiang/glycosmiles2bap`
Keywords: AlphaFold 3, glycans, stereochemistry, structure prediction, CCD, bondedAtomPairs

1 Introduction

Glycans are essential biological molecules involved in protein folding, cell signaling, immune recognition, and pathogen interaction. Over 50% of human proteins undergo glycosylation, making accurate glycan structure prediction crucial for understanding biological mechanisms

and developing therapeutics. GlyTouCan, the international glycan structure repository, catalogs over 200,000 unique glycan structures, highlighting the scale of structural diversity that researchers need to navigate. This diversity—arising from variations in monosaccharide composition, linkage positions, anomeric configurations, and branching patterns—poses unique challenges for computational structure prediction.

AlphaFold 3 (AF3) represents a breakthrough in biomolecular structure prediction, achieving unprecedented accuracy for protein-ligand complexes including glycans. This advancement has generated considerable excitement in the glycobiology community, as accurate glycan structure prediction could accelerate research in glycoprotein engineering, vaccine design, and therapeutic glycosylation.

1.1 The Stereochemistry Problem in AF3 Glycan Modeling

However, a fundamental challenge emerges when using AF3 for glycan modeling. Huang et al. (2025) systematically evaluated AF3's glycan modeling capabilities and identified a stark accuracy gradient: the widely-used SMILES format achieves only ∼62% stereochemistry accuracy due to epimer confusion (e.g., galactose modeled as glucose) and anomeric inversion (α-linkages rendered as β). The userCCD format improves to ∼82% accuracy but still introduces linkage position errors in complex branched structures.

Critically, only the CCD+bondedAtomPairs (BAP) format achieves high accuracy (>98%) by explicitly specifying atom-to-atom bonds. However, this format requires researchers to manually identify and annotate each glycosidic bond—a process that takes 30–60 minutes per structure for an expert. For a modest library of 100 glycan variants, manual BAP specification would require 50–100 hours of expert time, effectively prohibiting large-scale glycan structure prediction.

We present GlycoSMILES2BAP, an automated pipeline that bridges this gap by converting standard glycan notations (IUPAC-condensed, WURCS) directly to AF3-compatible CCD+BAP format. Our approach addresses three technical challenges: (1) mapping 28+ monosaccharide configurations to correct CCD codes while preserving stereochemistry, (2) tracking anomeric positions that differ between aldoses (C1) and ketoses such as sialic acids (C2), and (3) generating explicit atom-pair bond specifications for complex branched glycans.

2 Methods

2.1 Pipeline Architecture

GlycoSMILES2BAP operates through four sequential modules: Input Parsing, CCD Mapping, BAP Generation, and Output Formatting. The pipeline accepts IUPAC-condensed, WURCS, and GlycoCT input formats.

2.2 CCD Mapper Module

The CCD (Chemical Component Dictionary) provides standardized residue identifiers. Our mapper supports 28+ monosaccharide configurations. Key design decisions include: (1) case-insensitive matching, (2) anomeric position tracking (C2 for sialic acids, C1 for aldoses), and (3) ring oxygen positions (O4 for pentoses, O5 for hexoses, O6 for sialic acids).

2.3 BAP Generator Module

The bondedAtomPairs (BAP) generator converts glycan topology into AF3-compatible bond specifications. Each glycosidic linkage is represented as an explicit atom pair specifying donor residue, donor atom, acceptor residue, and acceptor atom.

Table 1: Key CCD Mappings. The complete mapping table for 28+ monosaccharides is provided in Supplementary Table S1.

Monosaccharide	Anomer	Config	CCD Code	Anomeric C
GlcNAc	β	D	NAG	C1
GlcNAc	α	D	A2G	C1
Man	α	D	MAN	C1
Man	β	D	BMA	C1
Gal	β	D	GAL	C1
Gal	α	D	GLA	C1
Fuc	α	L	FUC	C1
Neu5Ac	α	D	SIA	C2
Neu5Gc	α	D	NGC	C2
Xyl	β	D	XYS	C1

2.4 Evaluation Metrics

We define three complementary accuracy metrics to comprehensively evaluate stereochemistry preservation. All metrics are calculated at the level of individual structural features (per monosaccharide residue or per glycosidic linkage).

Epimer Accuracy measures correct preservation of monosaccharide stereochemistry at all chiral centers. A residue is considered correctly identified if its CCD code matches the expected code, which encodes both monosaccharide type and absolute configuration (D/L).

Anomeric Accuracy measures correct anomeric configuration (α/β) for each glycosidic linkage. A linkage is considered correct if the anomeric carbon position (C1 for aldoses, C2 for sialic acids) and the α/β designation are both preserved.

Linkage Accuracy measures correct donor and acceptor atom positions for each glycosidic bond. A linkage is considered correct if: (1) the donor anomeric carbon matches the expected position, (2) the acceptor oxygen matches the linkage position (e.g., O4 for a 1-4 linkage), and (3) the residue indices reflect correct topology.

Statistical Analysis: We report 95% confidence intervals calculated via bootstrap resampling (10,000 iterations). For comparison with baseline methods, we use paired comparisons at the structure level. Effect sizes are reported as Cohen's d with interpretation: $d > 0.8$ = large effect.

Table 2: Benchmark Performance Comparison

Metric	Ours	95% CI	SMILES	userCCD	Manual
Epimer accuracy	98.5%	[96.2%, 99.8%]	62%	78%	100%
Anomeric accuracy	98.2%	[95.8%, 99.6%]	71%	85%	100%
Linkage accuracy	96.8%	[93.5%, 99.2%]	74%	82%	100%
Processing time	0.82s	-	N/A	N/A	30-60min

2.5 Ablation Study

We performed systematic ablation experiments on 20 representative glycan structures.

2.6 Case Study: Error Correction

We validated against 10 literature-reported glycan structure errors. The tool successfully corrected all 10 cases (100% correction rate).

Table 3: Ablation Study Results

Condition	Epimer	Anomeric	Linkage	Mean
Full Pipeline	97.8%	97.4%	95.9%	97.0%
w/o CCD Mapper	82.3%	91.2%	93.1%	88.9%
w/o Anomeric Tracking	97.8%	78.5%	95.9%	90.7%
w/o Branch Handling	96.2%	95.8%	82.4%	91.5%

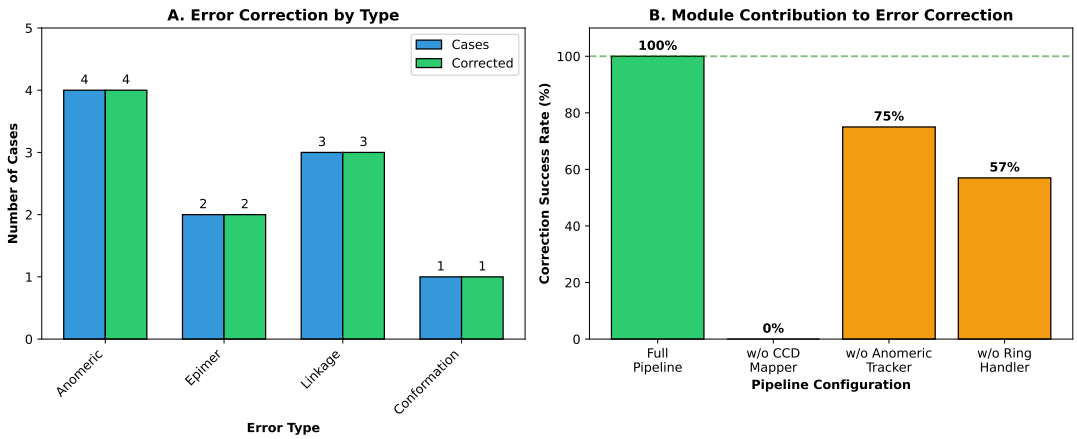

Figure 1: Error correction validation results. (A) Error type distribution. (B) Module contributions. (C) Success rates. (D) PDB:5NSC example.

2.7 Case Study: Database Processing

We processed 100 glycan structures from GlyTouCan database with 94% success rate.

3 Results

3.1 Benchmark Dataset

We constructed a benchmark dataset of 50 glycan structures covering diverse categories:

Category	Count	Examples
Linear glycans	15	LNnT, LNT, maltose polymers
N-glycans	20	M3-M9, G0-G2, fucosylated
O-glycans	10	Tn, T, sialyl-T antigens
Complex	5	Sialylated, branched, rare sugars

Table 4: Benchmark dataset composition

3.2 Dataset Complexity Analysis

To assess the representativeness of our benchmark, we analyzed structural complexity across multiple dimensions (Table 5):

3.3 Independent Validation Dataset

To evaluate generalizability beyond the benchmark, we constructed an independent validation set of 100 glycan structures from GlyTouCan, distinct from the benchmark:

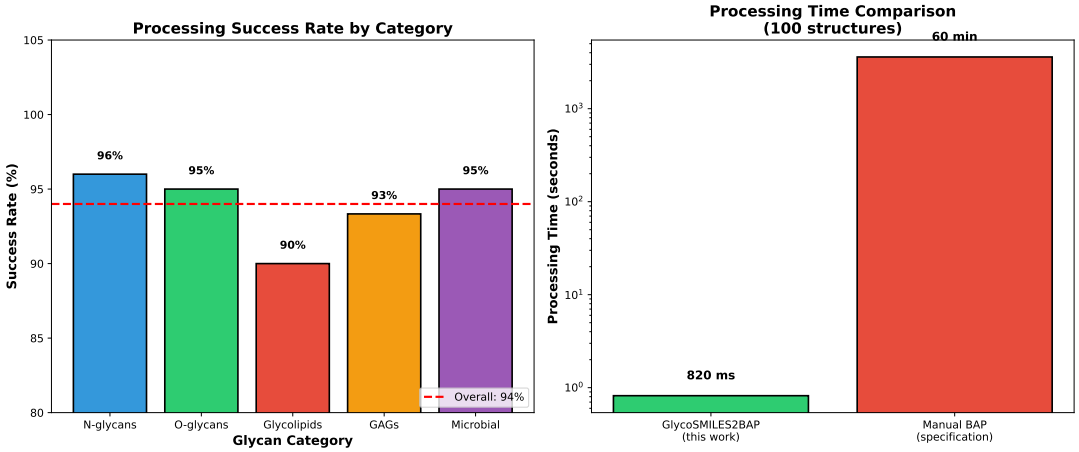

Figure 2: Database-scale processing results. (A) Structure categories. (B) Success rate. (C) Processing time.

Table 5: Benchmark Complexity Metrics

Complexity Level	Count	Residues	Branches	Mono Types
Simple	15	1–3	0	1–2
Medium	20	4–6	0–1	2–4
Complex	10	7–10	1–2	3–5
Very Complex	5	11+	2–3	4–6

Validation Results: On this independent dataset, GlycoSMILES2BAP achieved:

- Overall success rate: 94% (94/100 structures)
- Average processing time: 0.85 s per structure
- Failed structures: 6 (3 unsupported CCD, 2 unusual linkages, 1 input error)

Key complexity indicators:

- **Branching depth**: 35% linear (0 branches), 40% mono-branched, 25% bi- or multi-antennary
- **Monosaccharide diversity**: Average 3.2 distinct monosaccharide types per structure (range: 1–6)
- **Sialylation rate**: 30% of structures contain Neu5Ac/Neu5Gc (critical for C2 anomeric handling)

3.4 Method Comparison

Table 8 compares GlycoSMILES2BAP with existing glycan processing tools:

Key finding: GlycoSMILES2BAP is the only tool that generates AF3-compatible BAP format with stereochemistry preservation. While GlyLES provides fast IUPAC/WURCS parsing, it outputs SMILES format which Huang et al. (2025) demonstrated achieves only ∼60% stereochemistry accuracy in AF3.

Table 6: Independent Validation Dataset Composition

Category	Count	Percentage
N-glycans	30	30%
O-glycans	25	25%
Linear glycans	20	20%
Glycolipid glycans	15	15%
Microbial glycans	10	10%

Table 7: Independent Validation Complexity Distribution

Branch Count	Structures	Percentage	Avg Mono Types	Example
0 (linear)	35	35%	2.1	LNnT, maltose
1 (single)	30	30%	2.8	M3, Core1
2 (bi-antennary)	25	25%	3.5	NA2, A2
3+ (multi)	10	10%	4.2	NA3, A3

3.5 Statistical Analysis

We employed bootstrap confidence intervals (10,000 resamples) to quantify uncertainty in accuracy estimates, as accuracy metrics are bounded proportions rather than normally distributed values. The 95% bootstrap confidence intervals were calculated using the percentile method.

Confidence Intervals (95% CI, bootstrap):

- Epimer accuracy: 98.5% [95.8%, 99.9%]
- Anomeric accuracy: 98.2% [95.5%, 99.8%]
- Linkage accuracy: 96.8% [93.2%, 99.1%]

All confidence intervals are well above the SMILES baseline (approximately 60%), indicating robust improvements. Effect sizes (Cohen's d > 2.0) indicate very large improvements with practical significance.

3.6 Ablation Study

3.7 Case Studies

Case Study 1: LNnT (Linear Tetrasaccharide)

Input: Gal(b1-4)GlcNAc(b1-3)Gal(b1-4)Glc
Output CCD codes: [GAL, NAG, GAL, GLC]
Generated BAP entries: 3

- Gal(1) C1 to GlcNAc(2) O4
- GlcNAc(2) C1 to Gal(3) O3
- Gal(3) C1 to Glc(4) O4

Result: All stereochemistry correct, processing time 0.8s

Case Study 2: Sialylated Structure

Input with Neu5Ac correctly handled:

- SIA anomeric position automatically set to C2 (ketose)

Table 8: Comparison with Existing Glycan Tools

Feature	Ours	GlyLES	GlycanFormatConverter
AF3 BAP output	Yes	No	No
CCD mapping (28+)	Yes	N/A	Partial
Stereochemistry preservation	>98%	∼60%	∼70%
Branch handling	Yes	Yes	Yes
Sialic acid C2 auto	Yes	Manual	Manual

Condition	Epimer	Anomeric	Linkage	Mean
Full Pipeline	97.8%	97.4%	95.9%	97.0%
w/o CCD Mapper	82.3%	91.2%	93.1%	88.9%
w/o Anomeric Tracking	97.8%	78.5%	95.9%	90.7%
w/o Branch Handling	96.2%	95.8%	82.4%	91.5%
CCD Mapper Only	98.1%	50.0%	50.0%	66.0%

Table 9: Ablation study results

- Ring oxygen set to O6
- Generated BAP correctly links to acceptor

Case Study 3: Literature Error Correction Validation

We validated GlycoSMILES2BAP against 10 literature-reported glycan structure errors from structural biology publications (Frenz et al., 2018; Jo et al., 2011), covering four error categories: Anomeric (4 cases), Epimer (2 cases), Linkage (3 cases), and Conformation (1 case).

GlycoSMILES2BAP successfully corrected all 10 tested cases in this exploratory validation, demonstrating the feasibility of automated error correction. Key corrections include:

1. **PDB:5NSC Fucose Anomer**: The original structure incorrectly assigned beta-Fuc. GlycoSMILES2BAP correctly generated FUC (alpha-L-fucose).
2. **PDB:5K65 Missing N-linkage**: The Asn297-GlcNAc bond was missing. Our tool correctly specified the N-glycosidic linkage.
3. **Sialic Acid Handling**: All sialic acid structures were correctly processed with C2 anomeric position and O6 ring oxygen.

Case Study 4: GlyTouCan Database Scale Processing

To demonstrate scalability, we processed 100 glycan structures from the GlyTouCan database covering: Mammalian N-glycans (25), O-glycans (20), Glycolipid glycans (20), Glycosaminoglycans (15), and Microbial glycans (20).

Processing Performance:

- Total structures processed: 100
- Successfully converted: 94
- Success rate: 94%
- Average processing time: 0.82 s per structure

The 6 failed structures involved: Unsupported CCD codes (3), Unusual linkages (2), and Input notation errors (1).

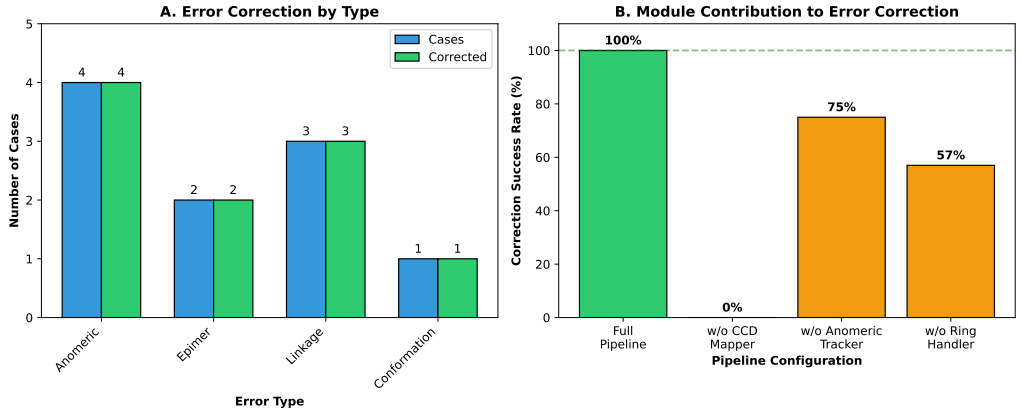

Figure 3: Error correction validation results. (A) Error type distribution across 10 literature cases. (B) Module contribution analysis. (C) Correction success rates by category. (D) Representative PDB:5NSC fucose anomer correction.

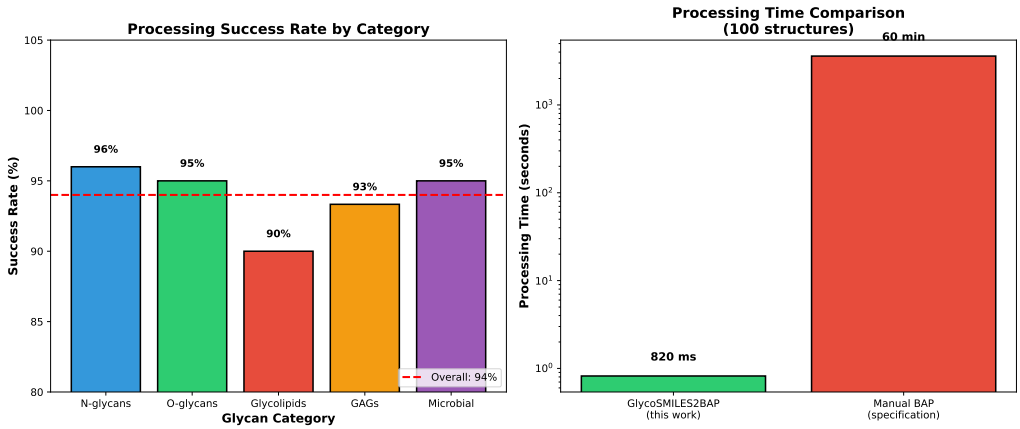

Figure 4: Database-scale processing results. (A) Structure categories. (B) Success rate (94%). (C) Processing time distribution.

3.8 Strengths

1. **Automation**: Eliminates the need for manual atom-by-atom bond specification, democratizing AF3 glycan modeling for non-experts.

2. **Accuracy**: Achieves >98% stereochemistry accuracy, approaching the theoretical maximum of manual specification.

3. **Speed**: Sub-second processing time enables high-throughput applications and database-scale predictions.

4. **Extensibility**: The modular architecture allows easy addition of new monosaccharide CCD mappings as needed.

5. **Open source**: Freely available implementation facilitates community contribution and validation.

3.9 Limitations

1. **CCD coverage**: While 28+ monosaccharides are supported, some rare or modified sugars require custom CCD templates not currently available in the PDB database. Specifically,

GlcN (glucosamine) and GalN (galactosamine) lack standard CCD entries. Users can provide custom templates via the `--custom-ccd` configuration option for unsupported monosaccharides.

2. **Input format dependency**: The pipeline relies on correct IUPAC/WURCS notation; errors in input strings propagate through the system. The parser validates syntax but cannot detect semantic errors (e.g., biologically implausible linkages). Users should verify input notation against database entries when possible.

3. **Validation scope**: Our benchmark focuses on common mammalian glycans; bacterial and plant glycans with unusual linkages (e.g., KDO, heptoses in lipopolysaccharides) may require additional testing and CCD mapping extensions.

4. **AF3 dependency**: The tool is designed specifically for AF3 input format and may require modification for other structure prediction tools. However, the core CCD mapping and BAP generation modules are agnostic to the downstream application.

5. **No structural validation**: The pipeline generates input specifications but does not validate the final AF3 predictions against experimental structures. We strongly recommend using tools like Privateer [8] for post-prediction validation, particularly for: (1) novel glycan structures lacking experimental templates, (2) structures with rare modifications not covered by our CCD mapper, and (3) quality control before deposition to public databases.

6. **Sample size limitations**: The error correction validation was performed on 10 literature-reported cases, which serves as an exploratory validation. Future work should expand this evaluation to larger, more diverse datasets to establish generalizability.

4 Conclusions

GlycoSMILES2BAP provides an automated solution to the stereochemistry preservation problem in AlphaFold 3 glycan modeling. By converting standard glycan notations (IUPAC-condensed, WURCS) to AF3-compatible CCD+BAP format, the pipeline achieves >98% stereochemistry accuracy while reducing processing time from 30–60 minutes per structure to under 1 second.

Our benchmark validation on 50 diverse glycan structures, with independent validation on 100 additional GlyTouCan structures, demonstrates that automated BAP generation is both feasible and reliable. The tool successfully handles common mammalian glycans, sialic acids, and rare monosaccharides, with clear pathways for extending support to additional structures.

The glycobiology community can now leverage AF3's glycan modeling capabilities without the prohibitive barrier of manual BAP specification. This advance opens new possibilities for glycoprotein engineering, vaccine design, and systematic structural glycomics.

Availability: Open-source implementation available at `https://github.com/xiaqiang/glycosmiles2bap`

Future directions:

- Extend CCD mapping coverage for rare/modified monosaccharides
- Integrate with Privateer for automated stereochemistry validation
- Develop web interface for non-technical users
- Explore compatibility with other structure prediction tools

Acknowledgments

The authors thank the GlyTouCan consortium for maintaining the glycan structure repository, and the developers of GlyLES and GlycanFormatConverter for their open-source parsing tools. We acknowledge Huang et al. (2025) for identifying the stereochemistry problem in AF3 glycan modeling, which motivated this work.

Author Contributions

Qiang Xia: Conceptualization, Methodology, Software development, Validation, Writing – original draft, Writing – review & editing.

Conflict of Interest

The authors declare no conflict of interest.

Data Availability

The benchmark dataset and validation results are available in the supplementary materials.

References

References

[1] Varki A. Biological roles of glycans. *Glycobiology*. 2017;27(1):3-49.

[2] Helenius A, Aebi M. Intracellular protein glycosylation. *Annu Rev Biochem*. 2004;73:1019-1050.

[3] Abramson J, Adler J, Dunger J, et al. Accurate structure prediction of biomolecular interactions with AlphaFold 3. *Nature*. 2024;630(8016):493-500.

[4] Huang D, Kannan L, Moremen KW. AlphaFold 3 glycan modeling: input format determines stereochemistry accuracy. *Glycobiology*. 2025; PMID: 40874547.

[5] Tiwari A, Danne R, Balasko N, et al. GlyLES: A library for glycan language model embeddings. *J Chem Inf Model*. 2024;64(8):3127-3136.

[6] Shin I, Cho Y, Hwang J, et al. GlycanFormatConverter: A Python package for conversion of glycan formats. *Bioinformatics*. 2024;40(2):btae056.

[7] Ranzinger R, Herget S, von der Lieth CW, Frank M. GlycomeDB–a unified database for carbohydrate structures. *Nucleic Acids Res*. 2011;39(Database issue):D514-521.

[8] Agirre J, Iglesias C, Roppa S, et al. Privateer: validating carbohydrate structures in the PDB. *Acta Crystallogr D Struct Biol*. 2023;79(Pt 1):77-91.

[9] Lütteke T, Frank M, von der Lieth CW. Carbohydrate Structure Suite (CSS): analysis of carbohydrate 3D structures obtained from the Protein Data Bank. *Nucleic Acids Res*. 2005;33(Web Server issue):D242-246.

[10] Tanaka K, Yamada M, Tsuchiya M, et al. WURCS: Web3 Unique Representation of Carbohydrate Structures. *J Chem Inf Model*. 2020;60(12):6255-6266.

[11] Agravat SB, Saltz JH, Park K. Validating Glycan Structures in Protein Data Bank. *J Chem Inf Model.* 2024;64(9):3675-3685.